

Hinglish Text Analysis: Challenges and Opportunities in Multilingual Natural Language Processing

Pragya Tewari^a, Sarthak Tyagi^b, Madhav Gumber^c, Paras Seekhwal^d

^apragyatewari@galgotiasuniversity.edu.in SCSE, Galgotias University, Greater Noida 203201, India

^bsarthaktyagi1704@gmail.com SCSE, Galgotias University, Greater Noida 203201, India

^cMadhav.gumber.12@gmail.com SCSE, Galgotias University, Greater Noida 203201, India

^dParasseekhwal@gmail.com SCSE, Galgotias University, Greater Noida 203201, India

Abstract: Hinglish is a blend of English and Hindi, commonly used in India, of late, is being used in general conversation exchanges in text messages, social media, and other digital platforms, especially in India. This designed mixture presents a challenge for language processing systems since Hinglish is unpredictable in structure, combines Roman and Devanagari scripts, and lacks standardized grammar. Traditionally, natural language processing systems have focused on single languages, making hybrid languages difficult to analyse. This paper proposes a specialized NLP pipeline to address these challenges in Hinglish. The pipeline includes preprocessing like script standardization, language identification, and tokenization, followed by a support vector machine and Naive Bayes classifier, and methods of deep learning using LSTM and BERT. Experimental results show that the proposed models outperform existing methods on correctly classifying Hinglish text and address unique complexities posed by code-mixed language systems. This in turn furthers the research in multilingual NLP, Supporting tool development in linguistically diverse settings..

1. Introduction

The use of Hinglish-a blend of Hindi and English language-as a way to communicate in social media and messaging applications, is fairly a reflection of the cultural and linguistic diversity of India. While this code-mixed language allows users to communicate more subtly and closer to reality, it nevertheless presents significant challenges for Natural Language Processing systems, which generally optimize single language datasets. The challenges towards the processing of Hinglish arise from its distinctive syntactic patterns, frequent switching within the same sentence between Hindi and English, and even lack of standardized grammar and spelling conventions. This makes tasks such as sentiment analysis, text classification, and language detection rather complex. (Das & Bandyopadhyay, 2010), which is actually culturally changing and digitizes communicative trends, presents additional challenges while developing reliable NLP models (Das & Bandyopadhyay, 2010)[1]. The present research has investigated the challenges that exist in the Hinglish text dimension and prepared novel schemes to address these weaknesses. We have tried to build these tools assisting in enhancing language processing of multilingual by developing better advanced preprocessing techniques, sentiment classification models, and performing a detailed error analysis.

2. Literature Review:

Natural Language Processing research has increasingly addressed unique challenges of multi-lingual and code-mixed languages in the recent past, driven by the growth of digital communication platforms and their limitations set by traditional monolingual models across the globe. Though Spanglish, Chinglish, and Arabic-English are among the most widely studied mixed languages, challenges concern Hinglish-one of the most prevalent hybrid languages-and distinctly differ from those mentioned. This blend of Hindi and English is mainly used to communicate informally through WhatsApp, Twitter, Facebook, and other social media platforms (Gupta & Choudhury, 2016) [2]. hence, Hinglish is an important candidate for sentiment analysis and social media monitoring. Several studies investigate the intricacies of Hinglish and other code-mixed languages. Chittaranjan et al. (2014) studied the code switching patterns in the Hinglish social media data and reported very frequent, unpredictable changes between Hindi and English in single sentences[3]. This study thus underlined the fact that these changes depend on the comfort factor of the user concerning the two languages and the context of the conversation. Sharma et al. extended the morphological challenges of Hinglish regarding tokenization and language identification. They discussed how the tokenization in Hinglish is really challenging because of mixed scripts, both Roman and Devanagari, along with informal spellings. The authors have proposed a hybrid approach which combines rule-based methods with machine learning for better accuracy in tokenization by inferring more requirements of specialized preprocessing techniques when the analysis is performed on code-mixed text The proposal in the work of Jhamtani et al. 2017 was sentiment analysis in Hinglish using traditional machine learning-based models[4]. However, these models generally perform poorly since large

annotated datasets of Hinglish are scarce and can difficultly process the Romanized Hindi and English together. Most of the SA models, which are designed for monolingual data, consider code-mixing as a knowledge gap; therefore, when applied to Hinglish text their performance is poor. Further, these models work on the formal structure of languages rather than the informal and grammatically inconsistent style typical of Hinglish. While recent works in multilingual models promise to improve the performance for Hinglish tasks, too, due to shared linguistic features between the languages, by virtue of studies such as Multilingual BERT by Devlin et al. (2019)[5], Hinglish presents challenges that are rather unique even to these models. The random switch-over to and fro between the languages in sentences, along with colloquial expressions and slang, demands more precise fine-tuning. Joshi et al. (2020) have done a deeper analysis of the code-mixed Hinglish data, incorporating several deep learning models, such as the LSTM and attention-based models[6]. They proposed new benchmark datasets for sentiment analysis in Hinglish and demonstrated that deep learning models strongly outperform traditional machine learning approaches. Yet again, however, their results point to significant progress that is yet to be achieved, especially those regarding domain specific vocabulary and sarcasm, which are abundant in the informal variety of Hinglish text. Patel and Kumar (2021) investigated transfer learning for code-mixed language analysis. [7]. They showed that models fine-tuned on code-mixed datasets outperform significantly better than models developed based on monolingual data alone. The important thing their study underlined was domain-specific training and fine-tuning, boosting the performance of NLP models for code-mixed languages like Hinglish. Despite the development in recent years, research in Hinglish sentiment analysis is restricted, and most works are done on small datasets or specific tasks like language detection and tokenization. More crucially, there is a need for the development of more and larger diverse corpora toward building real-time application NLP systems that should effectively handle the Hinglish language, particularly for sentiment analysis, NER, and machine translation. While multilingual NLP has indeed come a long way, especially in the contexts of languages such as Spanglish and Chinglish, Hinglish presents a set of challenges that are singular in themselves and require focused research. The development of NLP systems for Hinglish is yet to reach an appreciable state, with wide gaps in data, model generalization, and domain-specific performance (Faruqui & Dyer, 2014). This paper, therefore, attempts at bridging these gaps to further existing research into presenting a specialized pipeline for text analysis in Hinglish and, in turn, contribute toward the advancement of multilingual NLP systems.

3. Methodology

It involves describing, in detail, the different steps in data collection, preprocessing, model training, and the evaluation of this research to carry out the task and arrive at a powerful Hinglish sentiment analysis model that

will handle code-mixed text appropriately using some traditional and state-of-the-art NLP techniques.

3.1 . Data Collection

The data collected for this study is a diverse dataset of Hinglish texts collected from different social media websites and platforms like Twitter, Facebook, and WhatsApp by using a code-mixing approach. This also included chat logs from messaging platforms to capture information about informal conversations. The dataset comprises a mix of Hindi and English in both Devanagari and Roman scripts, reflecting typical code-mixing patterns found in Hinglish. Most of the data was extracted from existing corpora that are publicly available. The remaining portion was crawled using self-built web crawlers, designed to fetch Hinglish texts from the web. Every entry in the data is of tokenized sentence entries with language switching, which can be used to perform very fine-grained linguistic variability analysis in Hinglish.

To ensure a balanced representation of sentiment (positive, negative, neutral), the dataset was manually labelled by native Hinglish speakers.




Fig 1. Distribution of Sentiment Analysis

Human annotators were employed to provide sentiment labels, ensuring dataset reflects the nuances of Hinglish sentiment. This manual labelling was cross validated to achieve high inter-annotator agreement, thereby reducing any inherent subjectivity in sentiment classification

3.2 Preprocessing

Since code-mixed languages pose special challenges, preprocessing will be the most important preparatory step of the Hinglish text for sentiment analysis. The preprocessing pipeline basically consists of a few major steps: tokenization, normalization, and language detection.

3.2.1 Tokenization

Tokenizing Hinglish text faces several challenges: the presence of language switching within and at sentence boundaries. Some sentences contain words in both Devanagari and Roman scripts. In this paper, we adopt a hybrid approach to tokenization by combining rule-based methods with language-aware strategies. Apart from this, a bilingual tokenizer was developed which segmented the text based on its script identification and language boundary detection so that the code mixed text could be tokenized

accurately. Due care has also been given to slang, abbreviations, and hashtags that usually occur in Hinglish social media data. 3.2.2

3.2.2 Normalization

Normalization reduces the variability caused by non-standard spellings and mixed scripts prevalent in Hinglish text. We first normalized all the Devanagari-scripted Hindi words into Romanized versions by using a script conversion tool. In addition, we mapped informal spellings commonly encountered in Hinglish through a dictionary designed for spelling corrections, such as "kaise" to "kaisee". This dictionary was generated by merging human-submitted examples with machine suggestions identified through phonetic similarity. Normalization of the text cleaned up the dataset and reduced noise, which made for better tokenization and better consistency in downstream tasks.

3.2.3 Language Detection

After tokenization was performed and text normalization was done, the estimation of the language for each token was carried out. A language detection model was used that combined the strengths of character n-grams with lexical cues; this model classified each token as either Hindi or English. Since codeswitching is very frequent in Hinglish, this step becomes very important to carry out accurate text classification. The identification of the token's language is done to thereby enable treatments of Hindi and English tokens separately in higher analysis.

3.3 Sentiment Analysis

In this work, we explored two methods for sentiment classification: a traditional machine learning approach and a deep learning approach. Both are proposed as being well-suited for the code-mixed Hinglish text with high performance and good generalization over different domains.

3.3.1 Traditional Machine Learning Models

We started with some traditional models in machine learning like SVM and Naive Bayes because of their efficiency and simplicity on text classification problems. Features extracted from preprocessed Hinglish text include word-level n-grams, TF-IDF, and word embeddings (Conneau et al., 2017)[4]. These models are trained on the labelled dataset of Hinglish and evaluated by looking at a typical train-test split of the data. Although simple, these models were very useful to compare the performance of other more well-known deep learning approaches.

SVM

Support Vector Machine or SVM provides a far more robust approach toward predictive maintenance, as most data are not linearly separable [5] [3]. In finding an optimal hyperplane for the classification of data within a high-dimensional space, it is very effective in determining feature deviations or anomalies related to characteristics such as rotational speed and process temperature [10]. It uses kernel functions that change the data into higher dimensions for better classification. However, computationally expensive for large datasets and requiring a careful tuning of parameters such as the type of kernel and regularization in order to work optimally.


The equation for linear hyperplane:

language detection model was trained on a labelled Hinglish dataset, with both Hindi and English corpora serving as references. In this way, the model was able to deal effectively with tokens that are ambiguous between the two languages, given a context. It makes sure that accurate

Fig 2 . Support vector machine

Naïve Bayes

Support Vector Machine




The distance between data point x_i and decision boundary:

$$d_i = \frac{w^T x_i + b}{\|w\|} \quad (2)$$

Naive Bayes classifiers are a family of probabilistic models that use Bayes' theorem under a 'naive' assumption: Bayes assumes conditional independence of features, given the class label. Despite this so-called 'naive' assumption, they are widely used and very effective in many real-world applications, like text classification or spam filtering. They scale especially well with the number of features, requiring only a number of parameters proportional to the number of features, and can be fit using the highly efficient method of maximum likelihood estimation. However, their performance can degrade when the independence assumption is strongly violated.

$$P(A|B) = P(B|A)P(A) / P(B)$$



Naive Bayes

Fig 3. Naïve Bayes

3.3.2 Deep Learning Models

In order to overcome various challenges that code-mixed languages present to us, we intended to increase the classification accuracies by utilizing deep learning models, such as LSTMs and BERT. We have used the LSTM

model because it can learn long-

Algorithm	Accuracy	Precision	Recall	F1-score
SVM	0.82	0.80	0.81	0.81
Naïve Bayes	0.79	0.78	0.77	0.78
LSTM	0.85	0.84	0.83	0.84
BERT	0.88	0.87	0.86	0.87

Table 1 - Summarized Results of all Model accuracy.

range dependencies in sequential data. The handling of text from this area may come with frequent.

Also, we have fine-tuned mBERT using transfer learning techniques on our dataset in Hinglish (Devlin et al., 2019) [5]. Pretrained on many languages, mBERT allowed us to make good use of the shared linguistic features between Hindi and English (Chen et al., 2018) [6]. Fine-tuning means

sentiment categories: positive, negative, and neutral. More importantly,

4. Evaluation

We then evaluated our models using standard metrics such as accuracy, precision, recall, and F1-score for classification. The metrics comprehensively provided the performance of each model in handling a given task of sentiment classification across different categories. Apart from overall accuracy, we focused on precision and recall so that the models perform well on both positive and negative classes of sentiment.

Moreover, we conducted a very thorough error analysis to identify common misclassification patterns. The analysis focused on generally difficult cases, such as sarcastic remarks, ambiguous sentiment, and out-of-vocabulary words. By observing these instances where performance was poor, much insight was obtained into what needed to be improved, besides tokenization refinement for slang specific to the domain or optimization for ambiguous tokens.

5. RESULTS

Our experiments demonstrate that models trained exclusively on code-mixed Hinglish text outperform standard monolingual models significantly for SA tasks. The success can be attributed to our models being able to handle the complexities of code-switching, informal syntax, and script variability in Hinglish.

Traditional machine learning models like support vector machines and Naive Bayes yielded decent results with adequate performance in the basic sentence classification. However, they suffered from the subtleties of

adjusting the weights of this model so that it performs better on our specific task of Hinglish sentiment classification (Conneau et al., 2018) [8]. We experimented with different kinds of hyperparameters, such as learning rates, batch sizes, and sequence lengths, in order to see the best performance of the model on the codemixed text.

On the other hand, deep learning models show a remarkable improvement in performance. LSTM architecture with an attention mechanism proved very useful in modeling contextual information across code-mixed sentences. This model now could capture long range dependencies and mixed language input better; hence, it attained the highest accuracy score of 85%. Attention helped the model in focusing on the critical tokens that improve the classification of sentiment-laden words in both Hindi and English.

The overall top-performing model in our experiments was the fine-tuned multilingual BERT. By leveraging transfer learning from a pre-trained multilingually pre-trained model, mBERT, high generalization across different datasets facilitated this approach as highly versatile for handling Hinglish. Compared with the more traditional machine learning models, the LSTM model achieved a maximum accuracy of around 88% on this dataset. Thus, the BERT model with transformer-based architecture does the job very well for complicated sentence structures with variability in sentence length.

Moreover, the ability of BERT to capture context from both languages simultaneously helped it catch many nuances in sentiment, which other models mostly miss. Talking about overall metrics, the BERT model was excellent be it precision, recall, or F1-score-which indicated that it identified not only the correct sentiment but consistently did so across all

informal text and code-switching, misclassifying most sentences which had a high amount of mixing between Hindi and English. For example, the Naive Bayes classifier performed badly because of the context in codeswitched text, thus giving poor precision and recall scores.

BERT showed better robustness when applied to new, unseen Hinglish data and thus could generalize well for other code-mix language tasks. Detailed error analysis thereafter showed that most misclassifications fell either within highly ambiguous or sarcastic texts, where sentiment was generally difficult to interpret for even human annotators. Another problem the model faced was out-of-vocabulary words, mainly those bound to dialects or regional slang. These points could be considered further improvements in future work.

6. STUDY BENEFITS

The study of Hinglish sentiment analysis holds immense societal relevance, especially in a linguistically rich country like India, where code-mixed languages form an integral part of daily communication. With digital interactions playing an increasing role in personal, professional, and social contexts, understanding and processing languages like Hinglish is crucial for creating inclusive technological solutions. This research work addresses these challenges and helps develop tools and systems that cater to the needs of millions who use code-mixed languages frequently.

Another important practical impact of the research study may also be the implementation of improved customer engagements and sentiment

analyses across social media platforms. The machine learning models developed for sentiment analysis here are empowered to draw actionable insights for businesses out of customer feedback in the Hinglish language. This has, in turn, allowed businesses to perceive their audience, upgrade their customer support systems, and come up with marketing strategies matching the requirement of multilingual marketing. These tools also serve to identify public sentiment, track trends, and detect emerging issues to support real-time, data-driven decision-making.

The outcome of this research could, therefore, have enormous applications in the improvement of conversational AI systems, such as chatbots and virtual assistants. Enabling these technologies to process Hinglish without friction would allow them to interact more naturally and proficiently with users who comfortably switch between Hindi and English. It improves the accessibility and usability of AI-powered systems, especially for people who feel more at ease with code-mixed languages than their monolingual counterparts.

Moreover, this research furthers the causes of multilingual education and digital literacy. Application of accurate tools for translation, named entity recognition, and sentiment classification in Hinglish would help bridge the gap for users who seek access to educational resources or government services online. This is especially crucial in rural and semi-urban areas where code-mixed languages are more prevalent, thus ensuring that technological progress is inclusive and benefits underserved communities.

The conclusion: It furthers an important objective of state-of-the-art NLP inclusion and cultural sensitiveness. Narrowed down to focusing on Hinglish, the desperate need for multilingual communicative support lays a base from which further progress stands to be considered in other mixed languages. With any luck, the insight they gain here allows them to defy linguistic barriers in advanced technology and reach a way wherein meaningful digital relations are made and fostered for individuals, businesses, communities, and anyone else.

7. CONCLUSION

This points to the need for specialized models relevant to code-mixed languages, such as Hinglish, that pose unique challenges to the traditional NLP system. Our results show that the models conventionally trained on monolingual datasets yield poor performance in sentiment classification with Hinglish text. However, the overall performance increased significantly using the models trained on codemixed datasets. Deep learning models, for example, LSTM and fine-tuned BERT, can automatically handle some of the challenges posed by Hinglish text in sentence structures, informal grammars, and frequent switching between the two languages. The fine-tuned multilingual BERT gave the best results that had a strong balance in accuracy, precision, recall, and F1-score. This use of transfer learning allowed the model to generalize over different types of Hinglish datasets, and this solution could turn out to be very promising for real-world applications related to social media sentiment analysis, chatbots, and conversational AI systems. In particular, it can handle the context in a parallel manner in Hindi and English, making it especially appropriate for multilingual environmental applications where code switching is generally practiced. This work contributes to the growing research domain of multilingual NLP and indicates a more

sophisticated processing way for codemixed languages. Our study shows that, although Hinglish is essentially informal and non-standard, state-of-the-art NLP techniques are competent in handling such a language variety (Das et al., 2020)[13]. The work opens new avenues for similar attempts in other code mixed languages, such as Spanglish or Chinglish.

However, there are certain limitations in our approach. Although this work uses a vast dataset, it could have been updated from more sources to add to the diverse Hinglish text from various regions and even dialects. The inclusion of more examples in script-switch, such as Devanagari and Roman scripts, would help in developing more robustness in the models. Finetuning the models for other linguistic tasks, such as NER or machine translation, would extend the frontier of code-mixed NLP.

Future work will consider expanding the dataset of Hinglish to more informal text sources and further refining the said models for other code-mixed languages. Another line of future work could be the use of even more advanced techniques, including the use of hierarchical attention networks, and hybrid models that include traditional linguistic features in a deep learning architecture. We believe that such a contribution will help create a much more inclusive and productive multilingual NLP system, by further developing and fine-tuning such models for code-mixed languages.

REFERENCES

Chittaranjan, G., Vyas, Y., & Bali, K. (2014). Code-mixing: A Challenge for Language Identification in the Indian Context. *Proceedings of ACL Workshop on Language Technologies for Indian Social Media*.

Sharma, P., & Agarwal, A. (2016). Tokenization and Language Identification in Code-mixed Hinglish Text. *International Journal of Computational Linguistics*.

Jhamtani, H., Bhattacharyya, P., & Majumder, P. (2017). Sentiment Analysis for Hinglish Code-Mixed Text. *Proceedings of the International Conference on Computational Linguistics*.

Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *NAACL-HLT*.

Joshi, A., Prabhu, M., & Bhattacharyya, P. (2020). Deep Learning for Hinglish Sentiment Analysis: A Benchmark Dataset and Approach. *Journal of Multilingual NLP*.

Patel, N., & Kumar, R. (2021). Transfer Learning for Code-Mixed Language Analysis: An Empirical Study. *Proceedings of the International Conference on Artificial Intelligence*.

- Gupta, P., & Choudhury, M. (2016). Analysis of Code-mixed Text Using Multilingual Embeddings. *International Journal of Computational Linguistics*.
- Chen, X., Sun, Y., Athiwaratkun, B., Cardie, C., and Weinberger, K. (2018). Adversarial deep averaging networks for cross-lingual sentiment classification. *Transactions of the Association for Computational Linguistics*, 6:557– 570.
- Conneau, A., Lample, G., Ranzato, M., Denoyer, L., and Jégou, H. (2017). Word translation without parallel data. *CoRR*, abs/1710.04087.
- Conneau, A., Rinott, R., Lample, G., Williams, A., Bowman, S. R., Schwenk, H., and Stoyanov, V. (2018). Xnli: Evaluating cross-lingual sentence representations. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, Brussels, Belgium. Association for Computational Linguistics.
- A. S. Jadon, M. Parmar and R. Agrawal, "Hinglish Sentiment Analysis: Deep Learning Models for Nuanced Sentiment Classification in Multilingual Digital Communication," *2024 2nd International Conference on Device Intelligence, Computing and Communication Technologies (DICCT)*, Dehradun, India, 2024, pp. 318-323, doi: 10.1109/DICCT61038.2024.10533057.
- P. Awatramani, R. Daware, H. Chouhan, A. Vaswani and S. Khedkar, "Sentiment Analysis of Mixed-Case Language using Natural Language Processing," *2021 Third International Conference on Inventive Research in Computing Applications (ICIRCA)*, Coimbatore, India, 2021, pp. 651 - 658, doi: 10.1109/ICIRCA51532.2021.9544554.
- Das, A. and Bandyopadhyay, S. (2010). SentiWordNet for Indian languages. In *Proceedings of the Eighth Workshop on Asian Language Resources*, pages 56–63, Beijing, China.
- Das, A., Chakraborty, T., Solorio, T., Gamba, B., Aguilar, G., Kar, S., Garrette, D., and Pykl, S. (2020). Semeval-2020 task 9: Sentimix: Sentiment analysis for codemixed social media text. In *Proceedings of the 14th International Workshop on Semantic Evaluation (SemEval-2020)*. Association for Computational Linguistics.
- Kumar, A., Mangotra, A., Ailawadi, A., Jain, R., Arora, M. (2024). Sentiment Analysis on Multilingual Data: Hinglish. In: Swaroop, A., Polkowski, Z., Correia, S.D., Virdee, B. (eds) *Proceedings of Data Analytics and Management. ICDAM 2023. Lecture Notes in Networks and Systems*, vol 786. Springer, Singapore. https://doi.org/10.1007/978-981-99-6547-2_47
- Faruqui, M. and Dyer, C. (2014). Improving word space word representations using multilingual correlation. In *Proceedings of the 14th Conference of the European*